

This classic dataset contains the prices and other attributes of almost 54,000 diamonds. Perform single neuron regression to predict the diamond price.

price price in US dollars (\\$326--\\$18,823)

carat weight of the diamond (0.2--5.01)

cut quality of the cut (Fair, Good, Very Good, Premium, Ideal)

color diamond colour, from J (worst) to D (best)

clarity a measurement of how clear the diamond is (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best))

x length in mm (0--10.74)

y width in mm (0--58.9)

z depth in mm (0--31.8)

depth total depth percentage =  $z / \text{mean}(x, y) = 2 * z / (x + y)$  (43--79)

table width of the top of the diamond relative to widest point (43--95)

In [20]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
import tensorflow as tf
from sklearn.model_selection import train_test_split
```

In [3]:

```
df=pd.read_csv("diamonds.csv")
```

In [4]:

df

Out[4]:

	Unnamed: 0	carat	cut	color	clarity	depth	table	price	x	y	z
0	1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
...	...	...	...	...	...	...	...	...	...	...	...
53935	53936	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50
53936	53937	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61
53937	53938	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56
53938	53939	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74
53939	53940	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64

53940 rows × 11 columns

In [5]:

df.drop(['Unnamed: 0'],axis=1,inplace=True)

In [6]:

df

Out[6]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
...	...	...	...	...	...	...	...	...	...	...
53935	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50
53936	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61
53937	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56
53938	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64

53940 rows × 10 columns

In [7]:

```
df.insert(0, 'ID', range(1, 1 + len(df)))  
df.set_index("ID",inplace=True)
```

In [8]:

df

Out[8]:

	carat	cut	color	clarity	depth	table	price	x	y	z
ID										
1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
...	...	...	...	...	...	...	...	...	...	...
53936	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50
53937	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61
53938	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56
53939	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74
53940	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64

53940 rows × 10 columns

In [9]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 53940 entries, 1 to 53940
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   carat       53940 non-null  float64
1   cut         53940 non-null  object
2   color       53940 non-null  object
3   clarity     53940 non-null  object
4   depth       53940 non-null  float64
5   table       53940 non-null  float64
6   price       53940 non-null  int64
7   x           53940 non-null  float64
8   y           53940 non-null  float64
9   z           53940 non-null  float64
dtypes: float64(6), int64(1), object(3)
memory usage: 4.5+ MB
```

In [10]:

df\_num=df.select\_dtypes(['int64','float64'])

In [11]:

df\_cat=df.select\_dtypes(['object'])

In [12]:

```
from scipy.stats import skew
```

In [13]:

```
# Skew = 3 * (Mean - Median) / Standard Deviation.  
for i in df_num:  
    print(i, skew(df_num[i]))  
  
#data skewness > -1 and < 1 that means data is normally distributed
```

```
carat 1.1166148681277797  
depth -0.08229173779627727  
table 0.7968736878796518  
price 1.6183502776053016  
x 0.3786658120772097  
y 2.4340990250113648  
z 1.5223802221853722
```

In [14]:

```
from sklearn.preprocessing import LabelEncoder
```

In [15]:

```
le=LabelEncoder()  
for col in df_cat:  
    df_cat[col]=le.fit_transform(df_cat[col])
```

In [16]:

```
df_cat.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 53940 entries, 1 to 53940  
Data columns (total 3 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   cut         53940 non-null  int32  
1   color       53940 non-null  int32  
2   clarity     53940 non-null  int32  
dtypes: int32(3)  
memory usage: 1.0 MB
```

In [17]:

```
df_cat
```

Out[17]:

	cut	color	clarity
ID			
1	2	1	3
2	3	1	2
3	1	1	4
4	3	5	5
5	1	6	3
...	...	...	...
53936	2	0	2
53937	1	0	2
53938	4	0	2
53939	3	4	3
53940	2	0	3

53940 rows × 3 columns

In [18]:

```
df_new=pd.merge(df_num,df_cat,on="ID")
```

In [19]:

```
df_new
```

Out[19]:

	carat	depth	table	price	x	y	z	cut	color	clarity
ID										
1	0.23	61.5	55.0	326	3.95	3.98	2.43	2	1	3
2	0.21	59.8	61.0	326	3.89	3.84	2.31	3	1	2
3	0.23	56.9	65.0	327	4.05	4.07	2.31	1	1	4
4	0.29	62.4	58.0	334	4.20	4.23	2.63	3	5	5
5	0.31	63.3	58.0	335	4.34	4.35	2.75	1	6	3
...	...	...	...	...	...	...	...	...	...	...
53936	0.72	60.8	57.0	2757	5.75	5.76	3.50	2	0	2
53937	0.72	63.1	55.0	2757	5.69	5.75	3.61	1	0	2
53938	0.70	62.8	60.0	2757	5.66	5.68	3.56	4	0	2
53939	0.86	61.0	58.0	2757	6.15	6.12	3.74	3	4	3
53940	0.75	62.2	55.0	2757	5.83	5.87	3.64	2	0	3

53940 rows × 10 columns

In [21]:

```
x=df_new.drop("price",axis=1)
y=df_new["price"]
```

In [42]:

```
# Scaling the data
x_mean = np.mean(x)
```

In [43]:

```
x_mean
```

Out[43]:

```
carat      0.797940
depth      61.749405
table      57.457184
x          5.731157
y          5.734526
z          3.538734
cut        2.553003
color      2.594197
clarity    3.835150
dtype: float64
```

In [44]:

```
x_std = np.std(x)
```

In [45]:

```
x_std
```

Out[45]:

```
carat      0.474007
depth      1.432608
table      2.234470
x          1.121750
y          1.142124
z          0.705692
cut        1.027698
color      1.701089
clarity    1.724575
dtype: float64
```

In [48]:

```
x_x = (x-x_mean) / x_std
```



In [49]:

x\_x

Out[49]:

	carat	depth	table	x	y	z	cut	color	
ID									
1	-1.198168	-0.174092	-1.099672	-1.587837	-1.536196	-1.571129	-0.538099	-0.937163	-C
2	-1.240361	-1.360738	1.585529	-1.641325	-1.658774	-1.741175	0.434949	-0.937163	-1
3	-1.198168	-3.385019	3.375663	-1.498691	-1.457395	-1.741175	-1.511147	-0.937163	C
4	-1.071587	0.454133	0.242928	-1.364971	-1.317305	-1.287720	0.434949	1.414272	C
5	-1.029394	1.082358	0.242928	-1.240167	-1.212238	-1.117674	-1.511147	2.002131	-C
...	...	...	...	...	...	...	...	...	...
53936	-0.164427	-0.662711	-0.204605	0.016798	0.022304	-0.054888	-0.538099	-1.525021	-1
53937	-0.164427	0.942753	-1.099672	-0.036690	0.013548	0.100988	-1.511147	-1.525021	-1
53938	-0.206621	0.733344	1.137995	-0.063434	-0.047741	0.030135	1.407998	-1.525021	-1
53939	0.130927	-0.523105	0.242928	0.373383	0.337506	0.285204	0.434949	0.826413	-C
53940	-0.101137	0.314528	-1.099672	0.088115	0.118616	0.143499	-0.538099	-1.525021	-C

53940 rows × 9 columns

In [50]:

```
x_train, x_test, y_train, y_test = train_test_split(x_x, y, test_size=0.3)
```

In [51]:

```
model = tf.keras.Sequential([tf.keras.layers.Dense(1, input_shape=(9,))])
```

In [52]:

```
model.compile(optimizer="sgd", loss="mse")
```

In [53]:

```
trained_model = model.fit(x_train, y_train, epochs=50)
```

Epoch 1/50

1180/1180 [=====] - 1s 567us/step - loss: 4436778.4041

Epoch 2/50

1180/1180 [=====] - 1s 563us/step - loss: 1911444.9264

Epoch 3/50

1180/1180 [=====] - 1s 533us/step - loss: 1936684.9781

Epoch 4/50

1180/1180 [=====] - 1s 536us/step - loss: 1818138.3090

Epoch 5/50

1180/1180 [=====] - 1s 563us/step - loss: 1833496.5618

Epoch 6/50

1180/1180 [=====] - 1s 578us/step - loss: 1823795.9470

Epoch 7/50

1180/1180 [=====] - 1s 567us/step - loss: 1746740.7430

Epoch 8/50

1180/1180 [=====] - 1s 624us/step - loss: 1834848.8347

Epoch 9/50

1180/1180 [=====] - 1s 580us/step - loss: 1784011.5848

Epoch 10/50

1180/1180 [=====] - 1s 621us/step - loss: 1831464.1083

Epoch 11/50

1180/1180 [=====] - 1s 567us/step - loss: 1785456.6213

Epoch 12/50

1180/1180 [=====] - 1s 577us/step - loss: 1887061.7530

Epoch 13/50

1180/1180 [=====] - 1s 567us/step - loss: 1870181.1311

Epoch 14/50

1180/1180 [=====] - 1s 607us/step - loss: 1791370.8515

Epoch 15/50

1180/1180 [=====] - 1s 533us/step - loss: 1825887.4636

Epoch 16/50

1180/1180 [=====] - 1s 543us/step - loss: 1769436.5620

Epoch 17/50

1180/1180 [=====] - 1s 536us/step - loss: 1830362.8243

Epoch 18/50

1180/1180 [=====] - 1s 543us/step - loss: 1880613.8246

Epoch 19/50

1180/1180 [=====] - 1s 556us/step - loss: 1825285.1

722

Epoch 20/50

1180/1180 [=====] - 1s 543us/step - loss: 1807044.7

478

Epoch 21/50

1180/1180 [=====] - 1s 529us/step - loss: 1811000.4

961

Epoch 22/50

1180/1180 [=====] - 1s 584us/step - loss: 1812599.9

052

Epoch 23/50

1180/1180 [=====] - 1s 570us/step - loss: 1769915.1

559

Epoch 24/50

1180/1180 [=====] - 1s 550us/step - loss: 1856408.4

255

Epoch 25/50

1180/1180 [=====] - 1s 570us/step - loss: 1851703.5

313

Epoch 26/50

1180/1180 [=====] - 1s 577us/step - loss: 1791187.7

117

Epoch 27/50

1180/1180 [=====] - 1s 553us/step - loss: 1819278.1

862

Epoch 28/50

1180/1180 [=====] - 1s 536us/step - loss: 1739057.5

928

Epoch 29/50

1180/1180 [=====] - 1s 594us/step - loss: 1879389.5

278

Epoch 30/50

1180/1180 [=====] - 1s 612us/step - loss: 1761444.4

372

Epoch 31/50

1180/1180 [=====] - 1s 584us/step - loss: 1793397.4

790

Epoch 32/50

1180/1180 [=====] - 1s 624us/step - loss: 1845452.3

546

Epoch 33/50

1180/1180 [=====] - 1s 570us/step - loss: 1884488.8

819

Epoch 34/50

1180/1180 [=====] - 1s 604us/step - loss: 1814700.3

371

Epoch 35/50

1180/1180 [=====] - 1s 556us/step - loss: 1786090.4

794

Epoch 36/50

1180/1180 [=====] - 1s 539us/step - loss: 1860553.4

079

Epoch 37/50

1180/1180 [=====] - 1s 543us/step - loss: 1898835.5

630

Epoch 38/50

1180/1180 [=====] - 1s 587us/step - loss: 1852524.2

324

Epoch 39/50

1180/1180 [=====] - 1s 536us/step - loss: 1808252.7

326

```
Epoch 40/50
1180/1180 [=====] - 1s 621us/step - loss: 1815245.8
864
Epoch 41/50
1180/1180 [=====] - 1s 550us/step - loss: 1850231.7
570
Epoch 42/50
1180/1180 [=====] - 1s 539us/step - loss: 1886358.5
570
Epoch 43/50
1180/1180 [=====] - 1s 543us/step - loss: 1834787.6
779
Epoch 44/50
1180/1180 [=====] - 1s 536us/step - loss: 1760127.2
574
Epoch 45/50
1180/1180 [=====] - 1s 546us/step - loss: 1882329.5
883
Epoch 46/50
1180/1180 [=====] - 1s 546us/step - loss: 1900655.9
566
Epoch 47/50
1180/1180 [=====] - 1s 597us/step - loss: 1866726.5
197
Epoch 48/50
1180/1180 [=====] - 1s 621us/step - loss: 1755582.2
546
Epoch 49/50
1180/1180 [=====] - 1s 607us/step - loss: 1744955.9
602
Epoch 50/50
1180/1180 [=====] - 1s 590us/step - loss: 1828544.4
806
```

In [54]:

```
trained_model
```

Out[54]:

```
<tensorflow.python.keras.callbacks.History at 0x236b99fae50>
```

In [55]:

```
trained_model.history
```

Out[55]:

```
{'loss': [2686543.75,  
1915283.0,  
1865870.625,  
1846135.875,  
1831995.125,  
1833070.75,  
1830074.375,  
1829699.625,  
1828531.125,  
1829416.5,  
1827082.625,  
1830909.875,  
1830014.625,  
1829468.375,  
1830789.625,  
1829044.875,  
1829045.5,  
1829282.625,  
1828484.375,  
1829961.5,  
1828917.0,  
1827854.375,  
1829776.625,  
1831058.5,  
1827825.75,  
1827868.75,  
1828065.0,  
1829183.25,  
1828250.375,  
1831711.5,  
1828974.5,  
1828545.625,  
1829152.875,  
1828933.125,  
1828154.0,  
1828641.25,  
1829051.375,  
1829966.0,  
1828489.375,  
1830411.875,  
1829867.5,  
1830227.75,  
1827748.25,  
1827374.0,  
1829355.75,  
1828244.625,  
1831807.625,  
1832429.0,  
1825595.0,  
1828953.75]}}
```

In [56]:


```
trained_model.history['loss']
```

Out[56]:

```
[2686543.75,  
 1915283.0,  
 1865870.625,  
 1846135.875,  
 1831995.125,  
 1833070.75,  
 1830074.375,  
 1829699.625,  
 1828531.125,  
 1829416.5,  
 1827082.625,  
 1830909.875,  
 1830014.625,  
 1829468.375,  
 1830789.625,  
 1829044.875,  
 1829045.5,  
 1829282.625,  
 1828484.375,  
 1829961.5,  
 1828917.0,  
 1827854.375,  
 1829776.625,  
 1831058.5,  
 1827825.75,  
 1827868.75,  
 1828065.0,  
 1829183.25,  
 1828250.375,  
 1831711.5,  
 1828974.5,  
 1828545.625,  
 1829152.875,  
 1828933.125,  
 1828154.0,  
 1828641.25,  
 1829051.375,  
 1829966.0,  
 1828489.375,  
 1830411.875,  
 1829867.5,  
 1830227.75,  
 1827748.25,  
 1827374.0,  
 1829355.75,  
 1828244.625,  
 1831807.625,  
 1832429.0,  
 1825595.0,  
 1828953.75]
```

In [57]:

```
plt.plot(trained_model.history['loss'])  
plt.show()
```



In [58]:

```
y_hat = model.predict(x_test)
```

In [59]:

```
y_hat
```

Out[59]:

```
array([[ 461.8523 ],  
       [ 876.7219 ],  
       [6185.9927 ],  
       ...,  
       [ 5807.62   ],  
       [ 776.26196],  
       [-553.1023 ]], dtype=float32)
```

In [60]:

```
from sklearn.metrics import r2_score
```

In [61]:

```
r2_score(y_test, y_hat)
```

Out[61]:

```
0.8823440920404613
```

In [ ]:

